

STAT 231: Problem Set 2A

Kim Zhou

due by 5 PM on Monday, March 1

In order to most effectively digest the textbook chapter readings – and the new R commands each presents – series A homework assignments are designed to encourage you to read the textbook chapters actively and in line with the textbook’s Prop Tip of page 33:

“Pro Tip: If you want to learn how to use a particular command, we highly recommend running the example code on your own”

A more thorough reading and light practice of the textbook chapter prior to class allows us to dive quicker and deeper into the topics and commands during class. Furthermore, learning a programming language is like learning any other language – practice, practice, practice is the key to fluency. By having two assignments each week, I hope to encourage practice throughout the week. A little coding each day will take you a long way!

Series A assignments are intended to be completed individually. While most of our work in this class will be collaborative, it is important each individual completes the active readings. The problems should be straightforward based on the textbook readings, but if you have any questions, feel free to ask me!

Steps to proceed:

1. In RStudio, go to File > Open Project, navigate to the folder with the course-content repo, select the course-content project (course-content.Rproj), and click "Open"
2. Pull the course-content repo (e.g. using the blue-ish down arrow in the Git tab in upper right window)
3. Copy ps2A.Rmd from the course repo to your repo (see page 6 of the GitHub Classroom Guide for Stat231 if needed)
4. Close the course-content repo project in RStudio
5. Open YOUR repo project in RStudio
6. In the ps2A.Rmd file in YOUR repo, replace "YOUR NAME HERE" with your name
7. Add in your responses, committing and pushing to YOUR repo in appropriate places along the way
8. Run "Knit PDF"
9. Upload the pdf to Gradescope. Don’t forget to select which of your pages are associated with each problem. *You will not get credit for work on unassigned pages (e.g., if you only selected the first page but your solution spans two pages, you would lose points for any part on the second page that the grader can’t see).*

1. NYC Flights

a.

In Section 4.3.1, the `flights` and `carrier` tables within the `nycflights13` package are joined together. Recreate the `flightsJoined` dataset from page 80. Hint: make sure you've loaded the `nycflights13` package before referring to the data tables (see code on page 79).

```
library(nycflights13)
flights_joined <- flights %>%
  inner_join(airlines, by = c("carrier" = "carrier"))
```

b.

Now, create a new dataset `flightsJoined2` that:

- creates a new variable, `distance_km`, which is distance in kilometers (note that 1 mile is about 1.6 kilometers)
- keeps only the variables: `name`, `flight`, `arr_delay`, and `distance_km`
- keeps only observations where distance is less than 500 kilometers

Hint: see examples in Section 4.1 for subsetting datasets and creating new variables.

```
flights_joined2 <- flights_joined %>%
  mutate(
    distance_km = 1.6*distance
  ) %>%
  select(name, flight, arr_delay, distance_km) %>%
  filter(distance_km < 500)

head(flights_joined2)
```

```
## # A tibble: 6 x 4
##   name                flight arr_delay distance_km
##   <chr>              <int>     <dbl>      <dbl>
## 1 ExpressJet Airlines Inc.   5708      -14       366.
## 2 JetBlue Airways          1806       -4       299.
## 3 Southwest Airlines Co.   4646      -19       296
## 4 ExpressJet Airlines Inc.   4144       12       339.
## 5 JetBlue Airways          1002      -10       299.
## 6 JetBlue Airways           102        5       482.
```

c.

Lastly, using the functions introduced in Section 4.1.4, compute the number of flights (call this `N`), the average arrival delay (call this `avg_arr_delay`), and the average distance in kilometers (call this `avg_dist_km`) among these flights with distances less than 500 km (i.e. working off of `flightsJoined2`) *grouping by the carrier name*. Sort the results in descending order based on `avg_arr_delay`.

Getting NAs for `avg_arr_delay`? That happens when some observations are missing that data. Before grouping and summarizing, add a line to exclude observations with missing arrival delay information using `filter(is.na(arr_delay)==FALSE)`.

```

flights_joined2 %>%
  filter(is.na(arr_delay) == FALSE) %>%
  group_by(name) %>%
  summarize(
    num_flights = n(), avg_arr_delay = mean(arr_delay), avg_dist_km = mean(distance_km)
  ) %>%
  arrange(desc(avg_arr_delay))

```

```

## # A tibble: 11 x 4
##   name                num_flights avg_arr_delay avg_dist_km
##   <chr>                <int>         <dbl>         <dbl>
## 1 Mesa Airlines Inc.      286          18.0          360.
## 2 ExpressJet Airlines Inc. 14753         15.6          373.
## 3 Envoy Air              2741          11.0          351.
## 4 JetBlue Airways       13443           8.66          385.
## 5 Endeavor Air Inc.      6144           6.82          339.
## 6 Southwest Airlines Co.   200           4.92          272.
## 7 United Air Lines Inc.   3307           4.09          320.
## 8 SkyWest Airlines Inc.     1             3            366.
## 9 US Airways Inc.        9093           2.22          308.
## 10 American Airlines Inc.  1428           1.88          299.
## 11 Delta Air Lines Inc.   1201          -0.643          325.

```

2. Baby names

a.

Working with the `babynames` data table in the `babynames` package, create a dataset `babynames2` that only includes years 2000 to 2017.

```
library(babynames)

babynames2 <- babynames %>%
  filter(year %in% 2000:2017)
```

b.

Following the code presented in Section 5.2.4, create a dataset called `BabyNarrow` that summarizes the total number of people with each name (born between 2000 and 2017), grouped by sex. (Hint: follow the second code chunk on page 102, but don't filter on any particular names.) Look at the dataset. Why have we called this dataset "narrow"?

ANSWER:

```
BabyNarrow <- babynames2 %>%
  group_by(name, sex) %>%
  summarize(total = sum(n))
```

'summarise()' has grouped output by 'name'. You can override using the '.groups' argument.

BabyNarrow

```
## # A tibble: 73,332 x 3
## # Groups:   name [67,063]
##   name      sex  total
##   <chr>    <chr> <int>
## 1 Aaban      M      107
## 2 Aabha      F       35
## 3 Aabid      M       10
## 4 Aabir      M        5
## 5 Aabriella F       32
## 6 Aada       F        5
## 7 Aadam      M      202
## 8 Aadan      M      130
## 9 Aadarsh    M      199
## 10 Aaden     F        5
## # ... with 73,322 more rows
```

We call this dataset narrow because our categorical variable `sex` is combined into one column and the total counts are for each name, not grouped by sex, so we don't have as many columns.

c.

Now, following the code chunk presented on page 103*, put the data into a wide format (call the new dataset `BabyWide`), and only keep observations where both `M` and `F` are greater than 10,000. Compute the `ratio` (as `pmin(M/F, F/M)`) and identify the top three names with the largest ratio. (Note: these names could be different from the ones found on page 103 since we limited the dataset to years 2000-2017 and names with greater than 10,000 individuals.)

- Note: you can use the `pivot_wider()` function instead of the `spread()` function if using the 2nd edition of the textbook (e.g., see Section 6.2.2 and 6.2.3 in the 2nd edition). I find `pivot_wider()` and `pivot_longer()` to be more intuitive than `spread()` and `gather()`.

ANSWER:

```
BabyWide <- babynames2 %>%
  group_by(sex, name) %>%
  summarize(total = sum(n)) %>%
  spread(key = sex, value = total, fill = 0) %>%
  filter(M > 10000, F > 10000) %>%
  mutate(ratio = pmin(M/F, F/M)) %>%
  arrange(desc(ratio))
```

'summarise()' has grouped output by 'sex'. You can override using the '.groups' argument.

```
BabyWide%>%
  head(3)
```

```
## # A tibble: 3 x 4
##   name      F      M ratio
##   <chr> <dbl> <dbl> <dbl>
## 1 Justice 10947 11267 0.972
## 2 Skyler  17120 22154 0.773
## 3 Quinn  25022 19080 0.763
```

The top three balanced gender neutral names between 2000-2017 are Justice, Skyler, and Quinn.

d.

Lastly, use the `gather()` function (or the `pivot_longer()` function) to put the dataset back into narrow form. Call this dataset `BabyNarrow2`. Hint: see Section 5.2.3. Why are the number of observations in `BabyNarrow2` different from that in `BabyNarrow`?

ANSWER:

```
BabyNarrow2 <- BabyWide %>%
  gather(key = sex, value = total, F, M)

BabyNarrow2
```

```
## # A tibble: 50 x 4
##   name      ratio sex  total
##   <chr>    <dbl> <chr> <dbl>
## 1 Justice 0.972 F    10947
## 2 Skyler  0.773 F    17120
## 3 Quinn   0.763 F    25022
## 4 Amari   0.751 F    11778
## 5 Casey   0.720 F    12109
## 6 Riley   0.666 F    89827
## 7 Peyton  0.641 F    61217
## 8 Emerson 0.632 F    18592
## 9 Charlie 0.624 F    13255
## 10 Dakota 0.612 F    21950
## # ... with 40 more rows
```

The number of observations in BabyNarrow2 is much fewer than the number in BabyNarrow because BabyNarrow is the entire babynames2 data set. In contrast, BabyNarrow2 is made from BabyWide, which we had filtered down to the names with more than 10000 males and 10000 females. So since BabyWide only had 25 observations, BabyNarrow2 has 50 observations.